

Creating PyTorch Solutions on Azure



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com



Overview

Azure Machine Learning Service

Azure notebooks for PyTorch

PyTorch estimators on Azure

**Leverage Horovod for distributed
deep learning**


Azure Deep Learning VMs



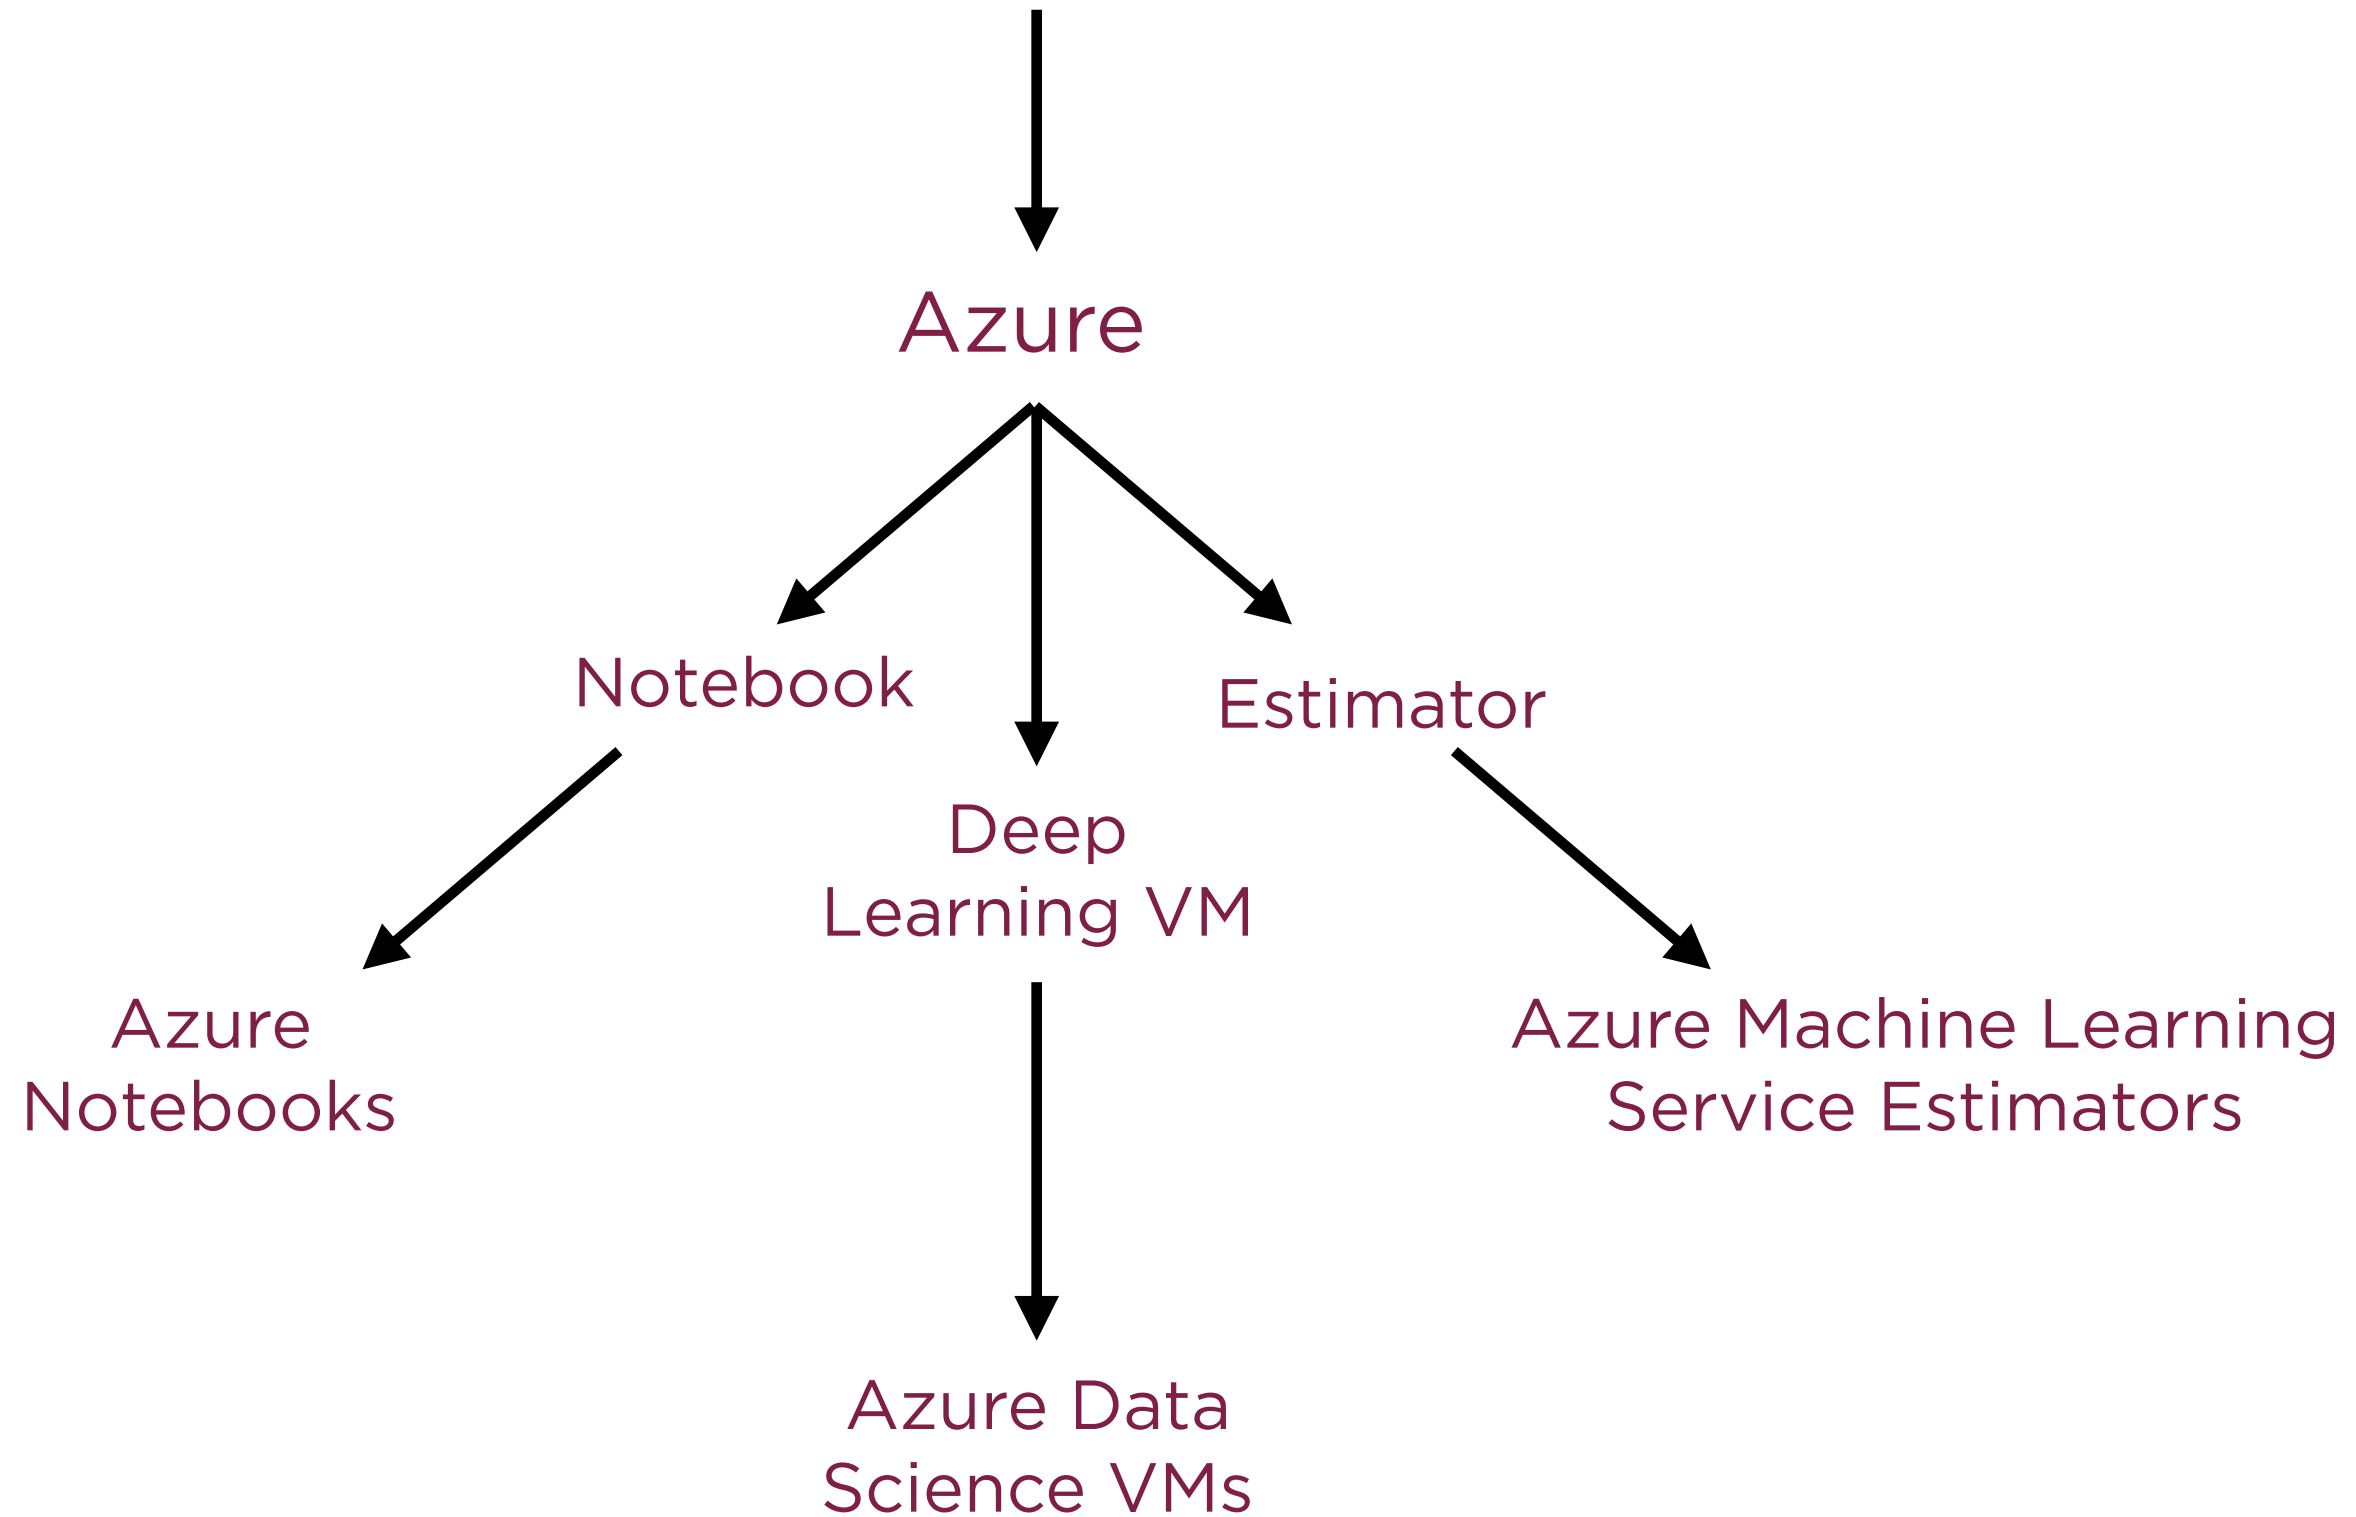
Machine Learning in Azure



PyTorch on the Cloud



PyTorch on the Cloud



Azure Machine Learning Service



Cloud service to train, deploy, automate and manage machine learning models

Supports open-source technologies such as PyTorch, TensorFlow and scikit-learn



Azure Machine Learning Service



Build custom models using open-source frameworks

Run distributed training

Perform hyperparameter tuning

Deploy models

Use models for prediction



Demo

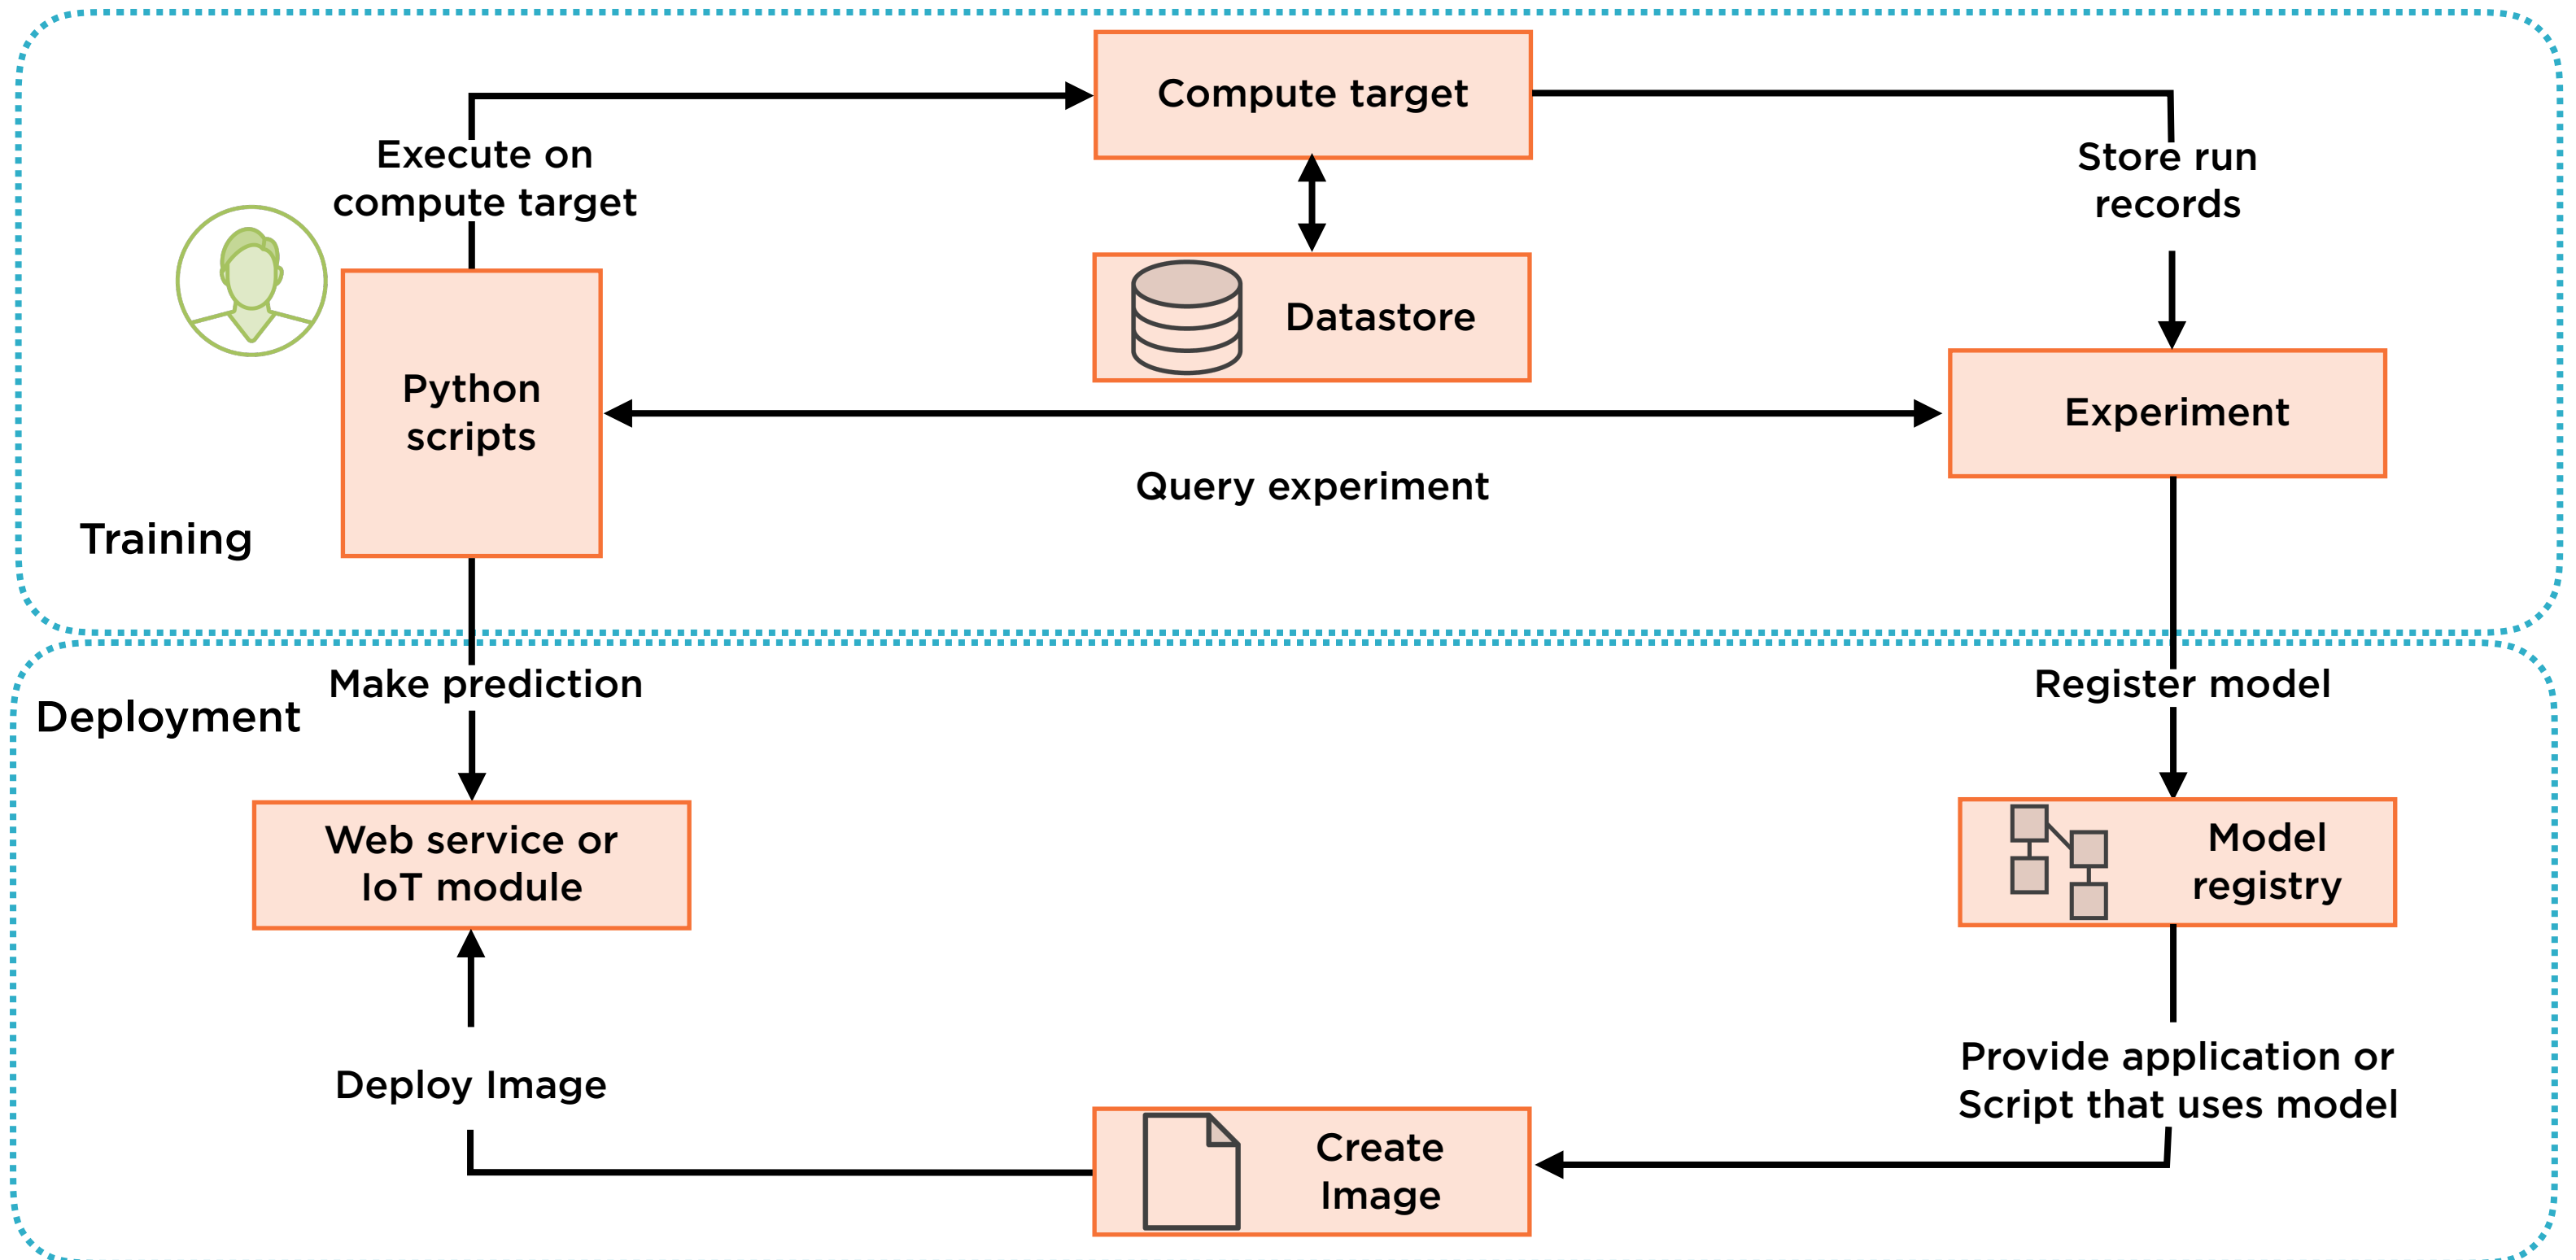
Prototyping PyTorch models on Azure notebooks



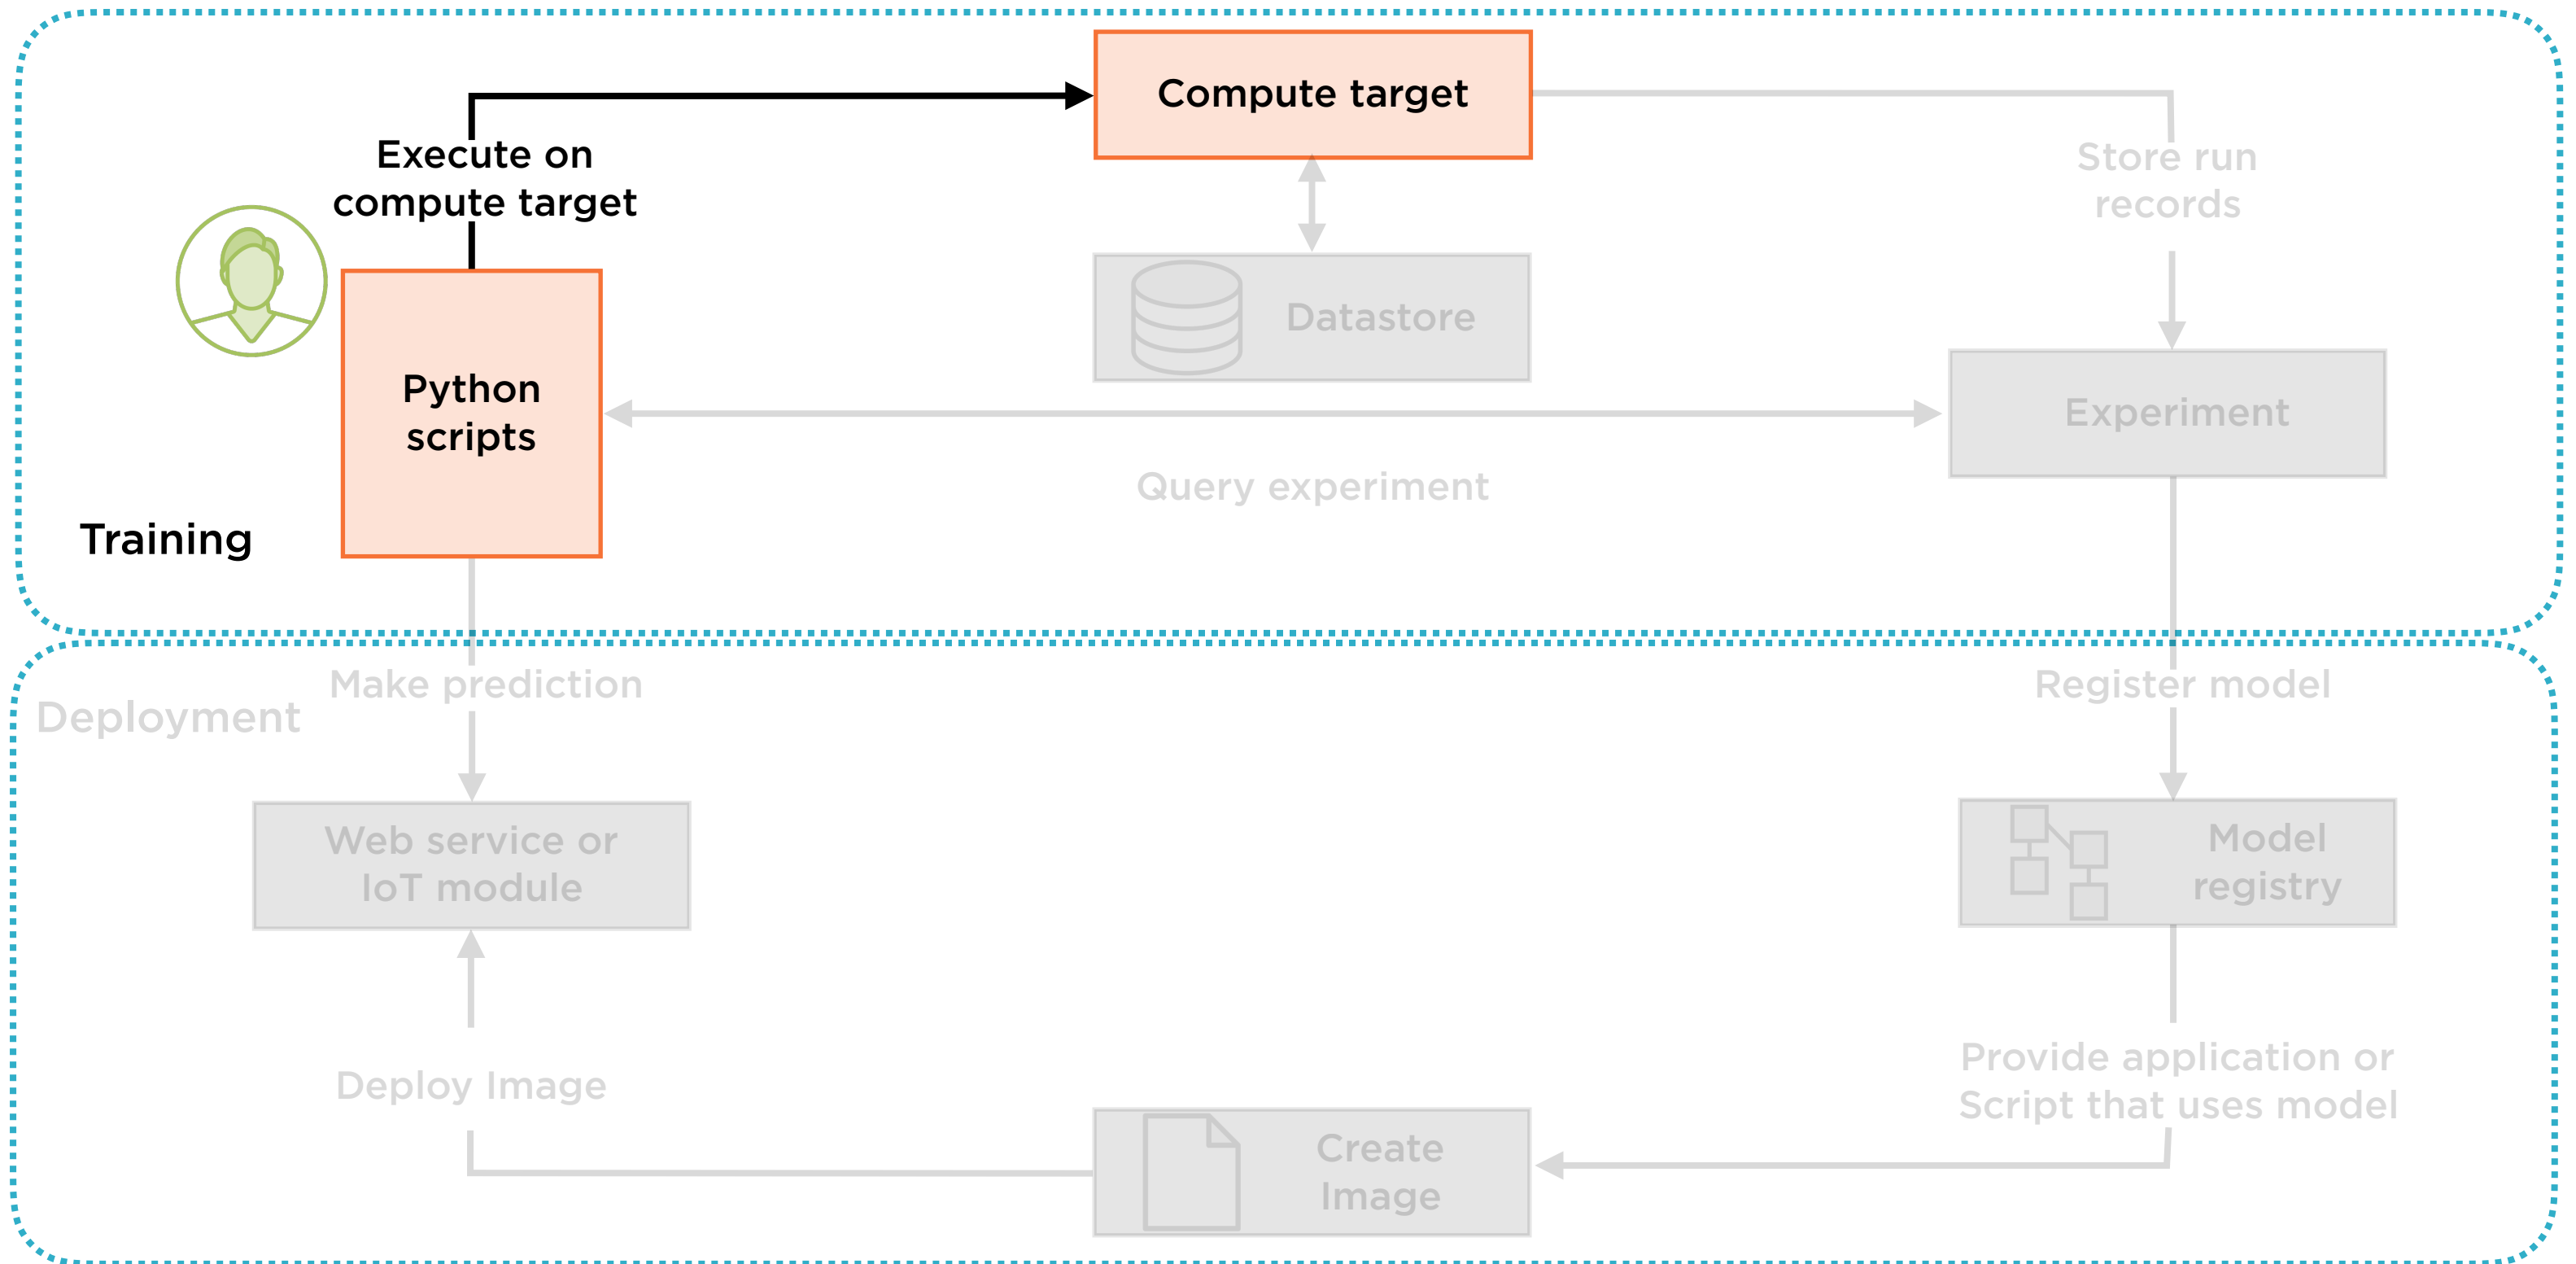
Azure Machine Learning Service



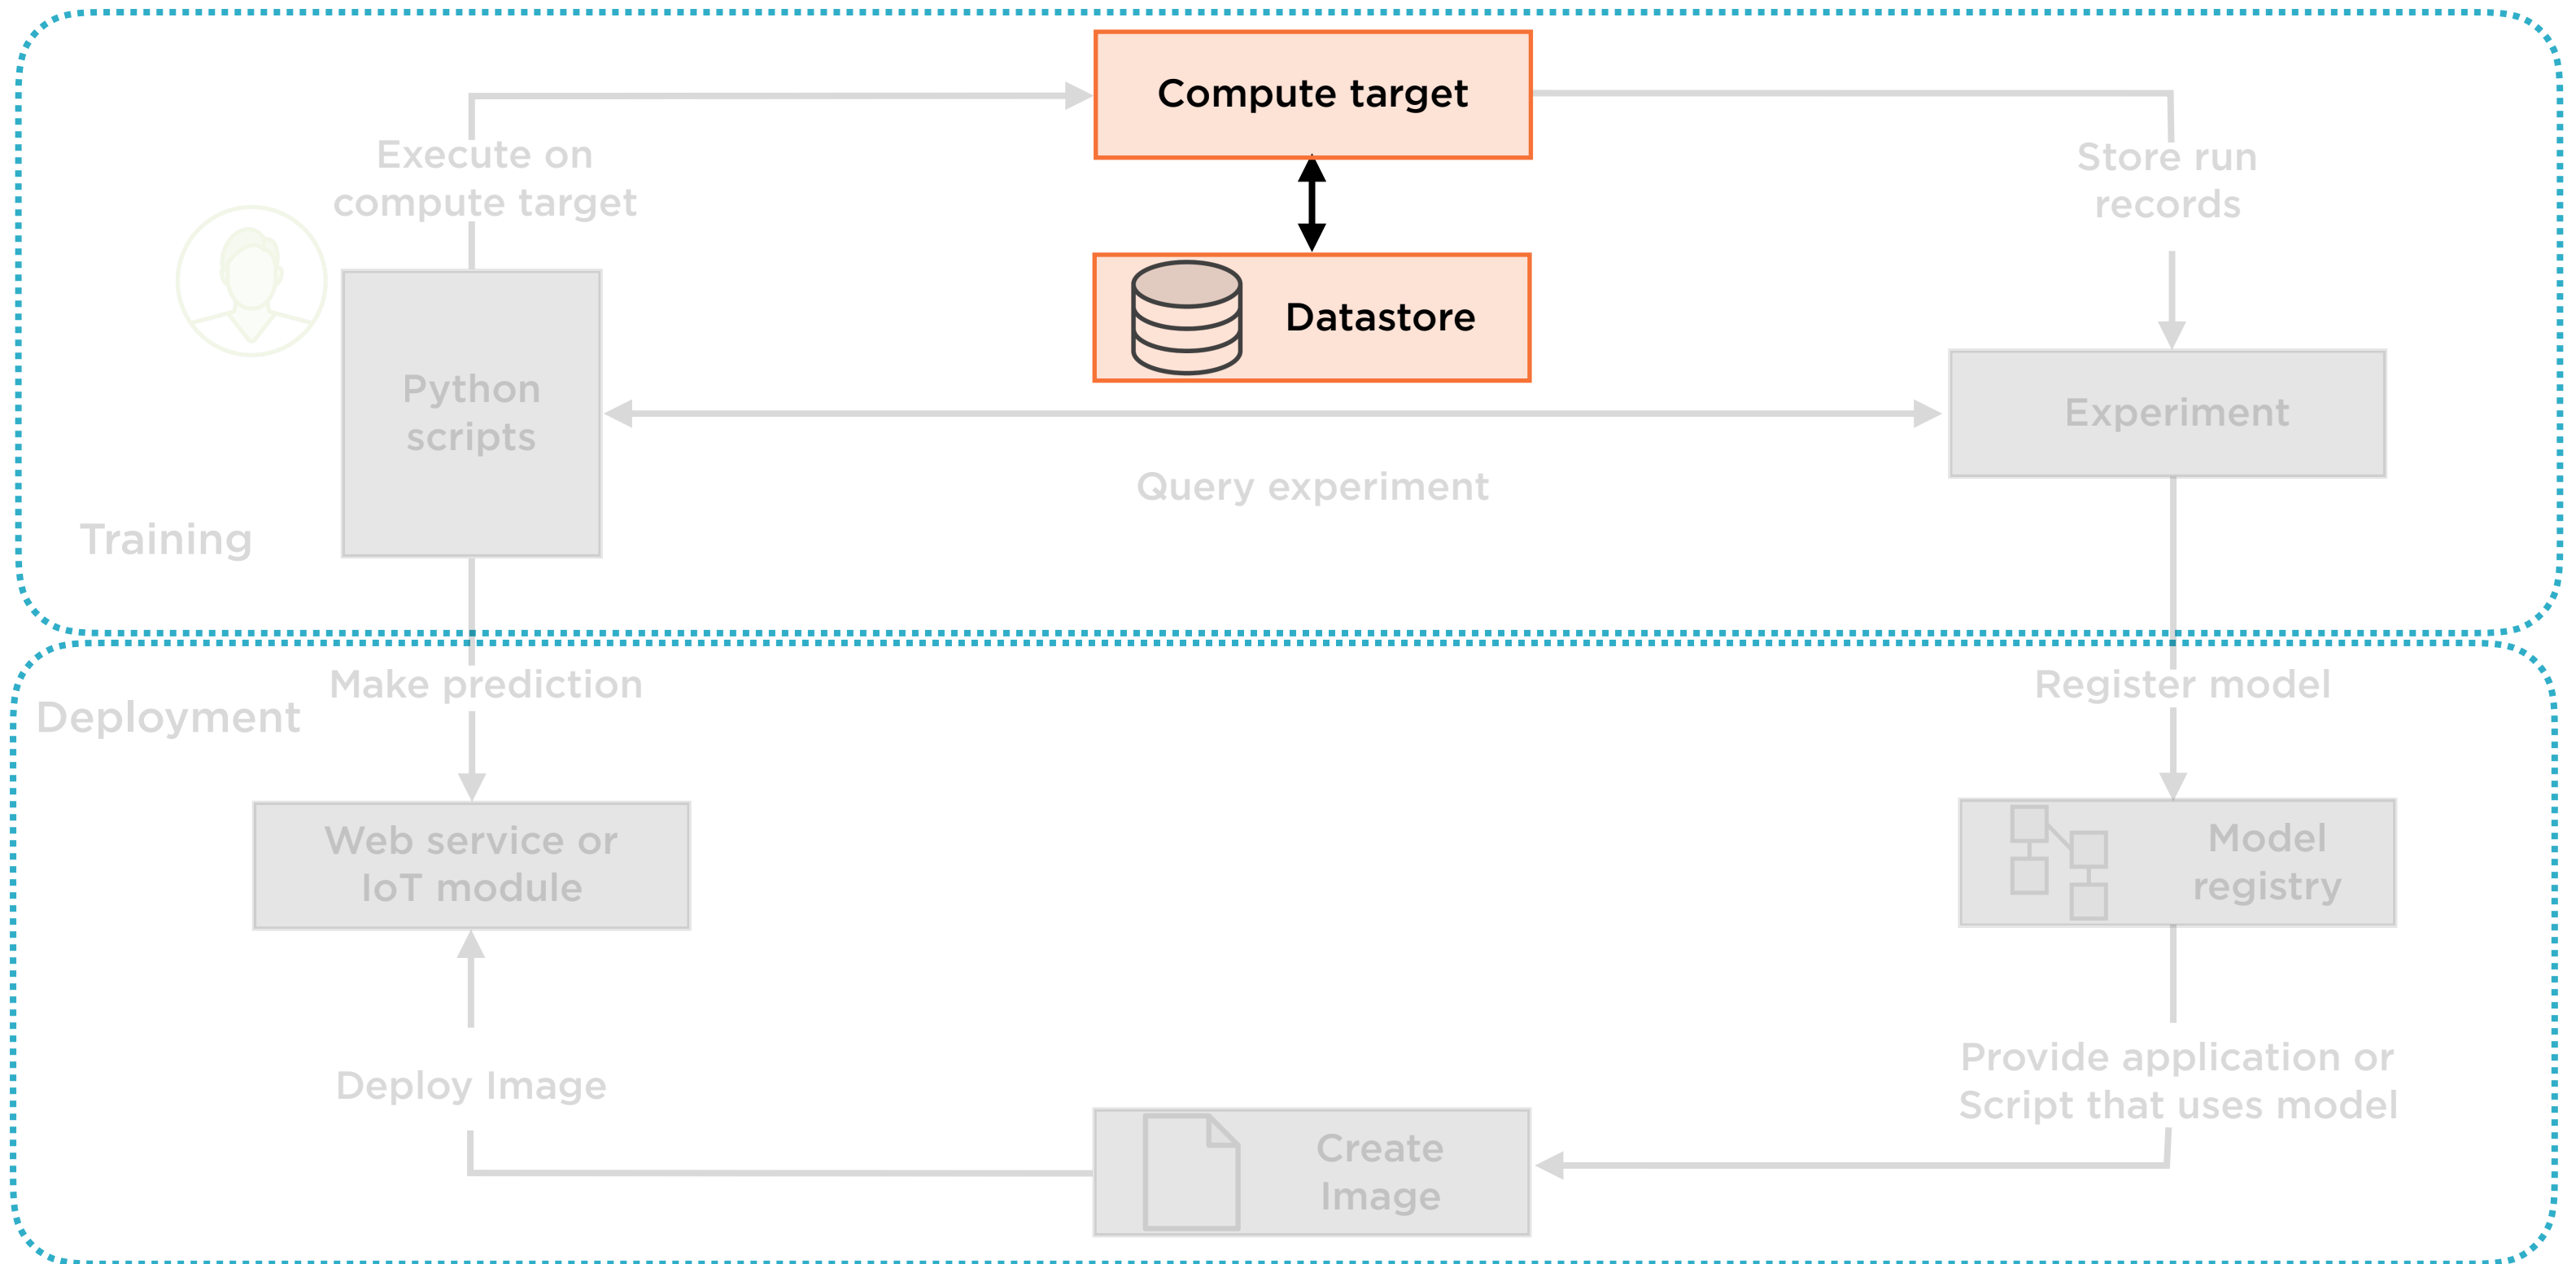
Azure Machine Learning Service



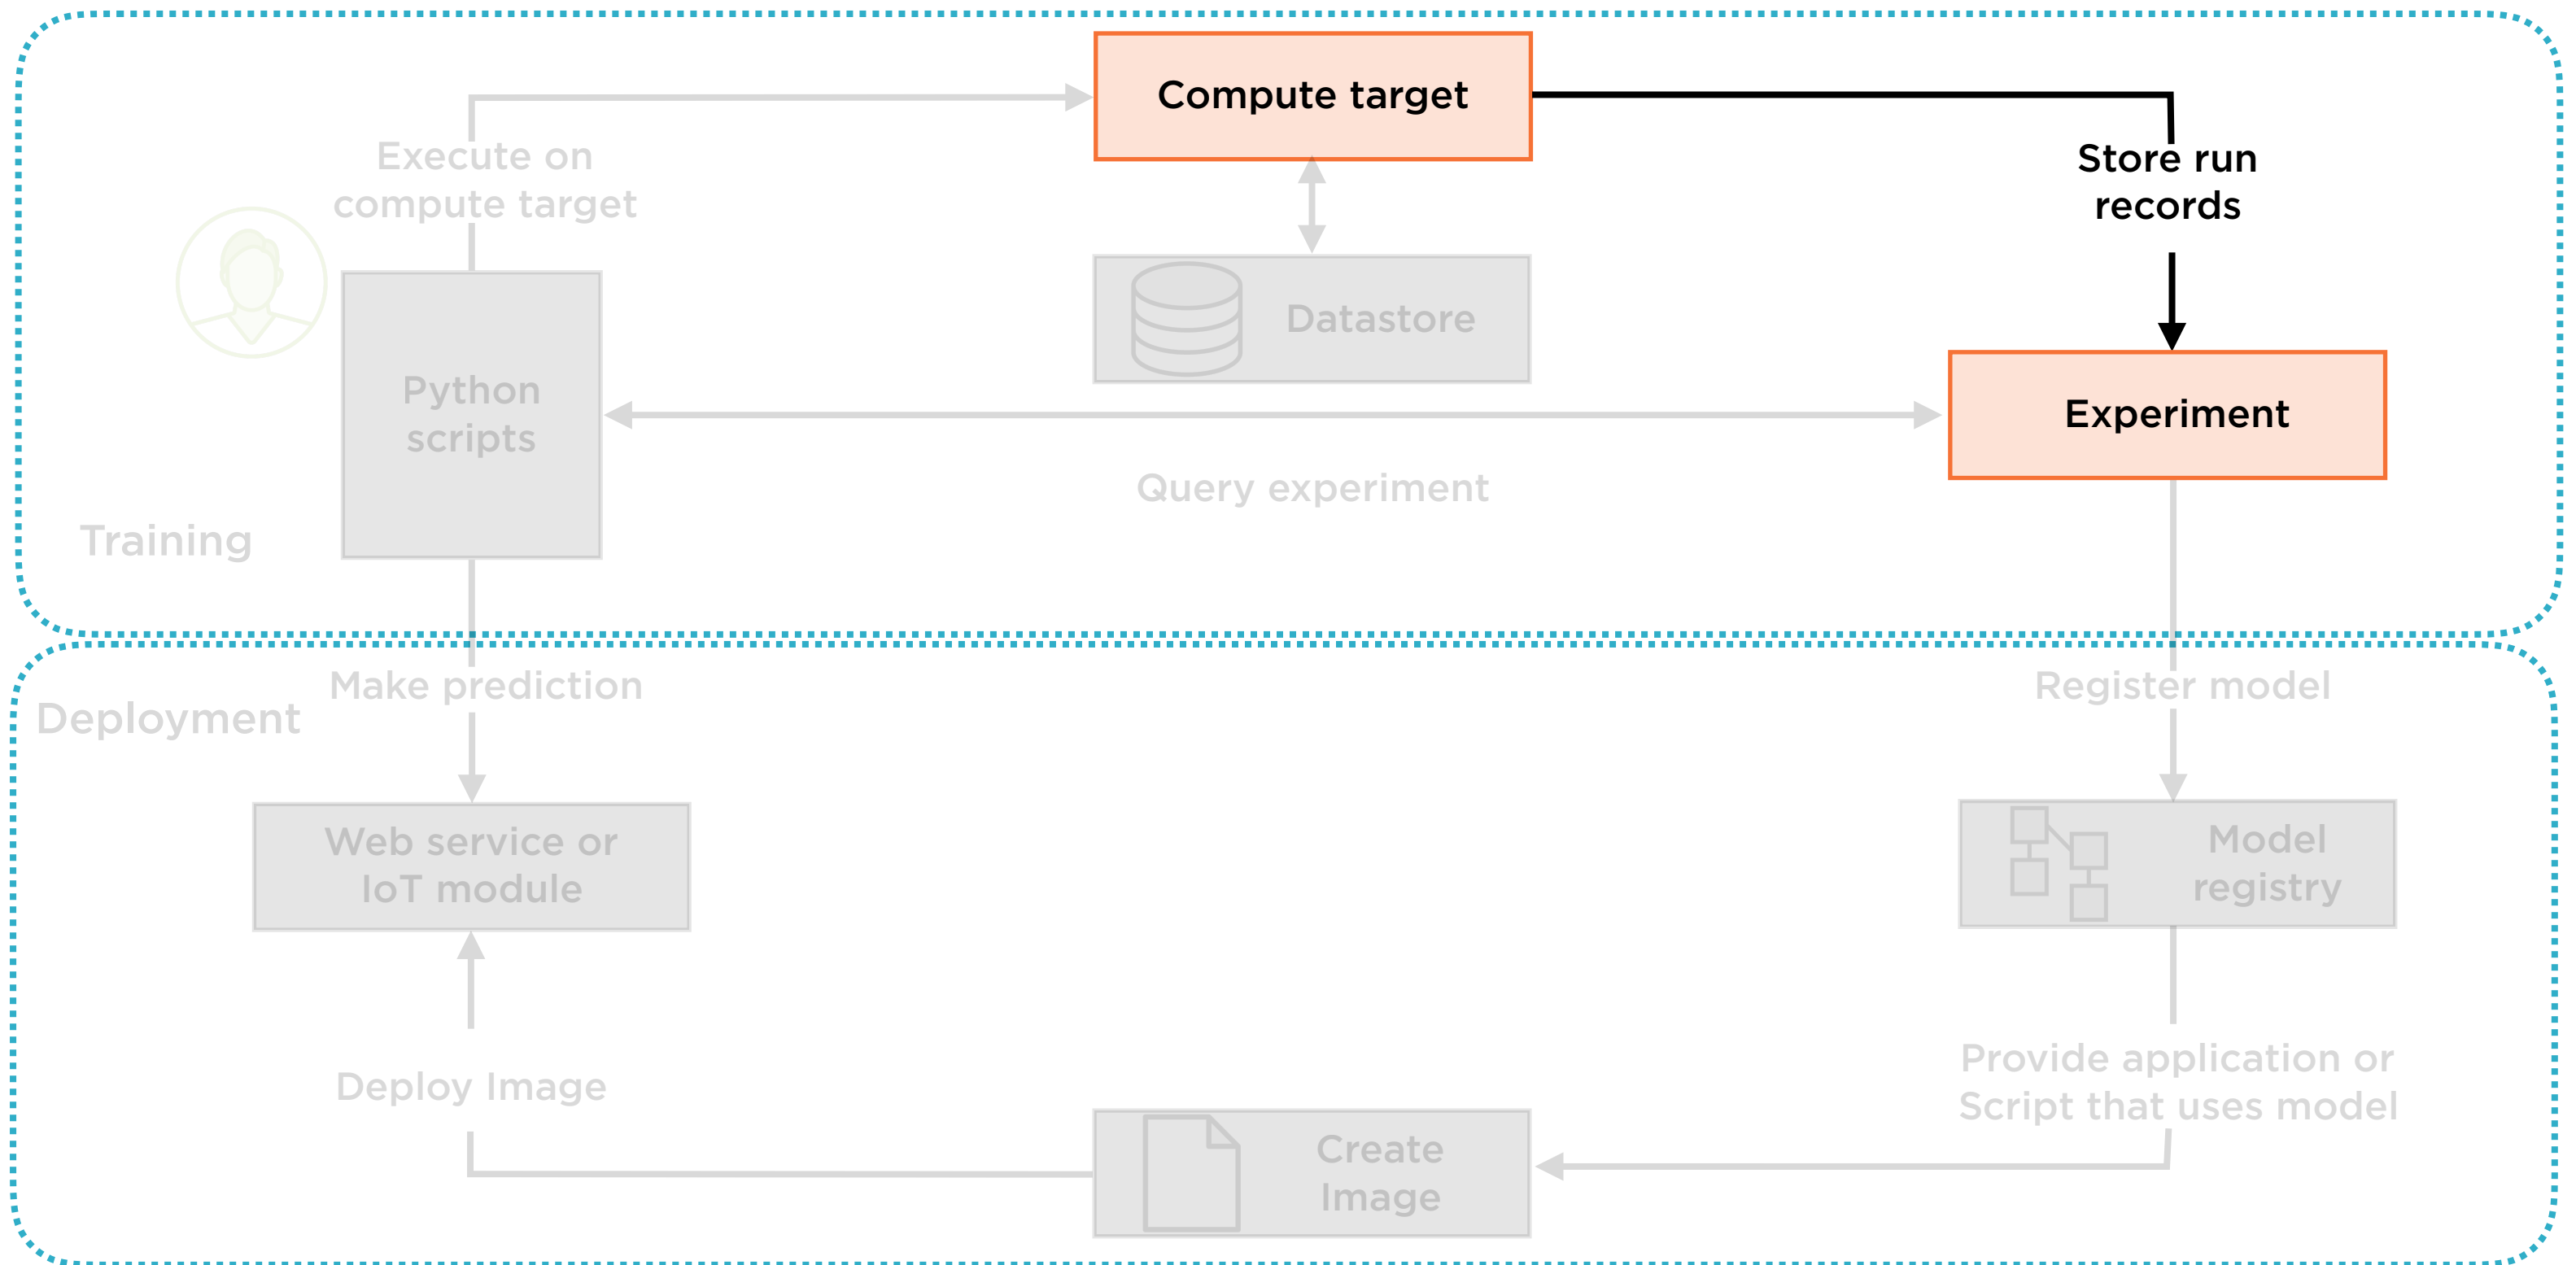
Compute Resources



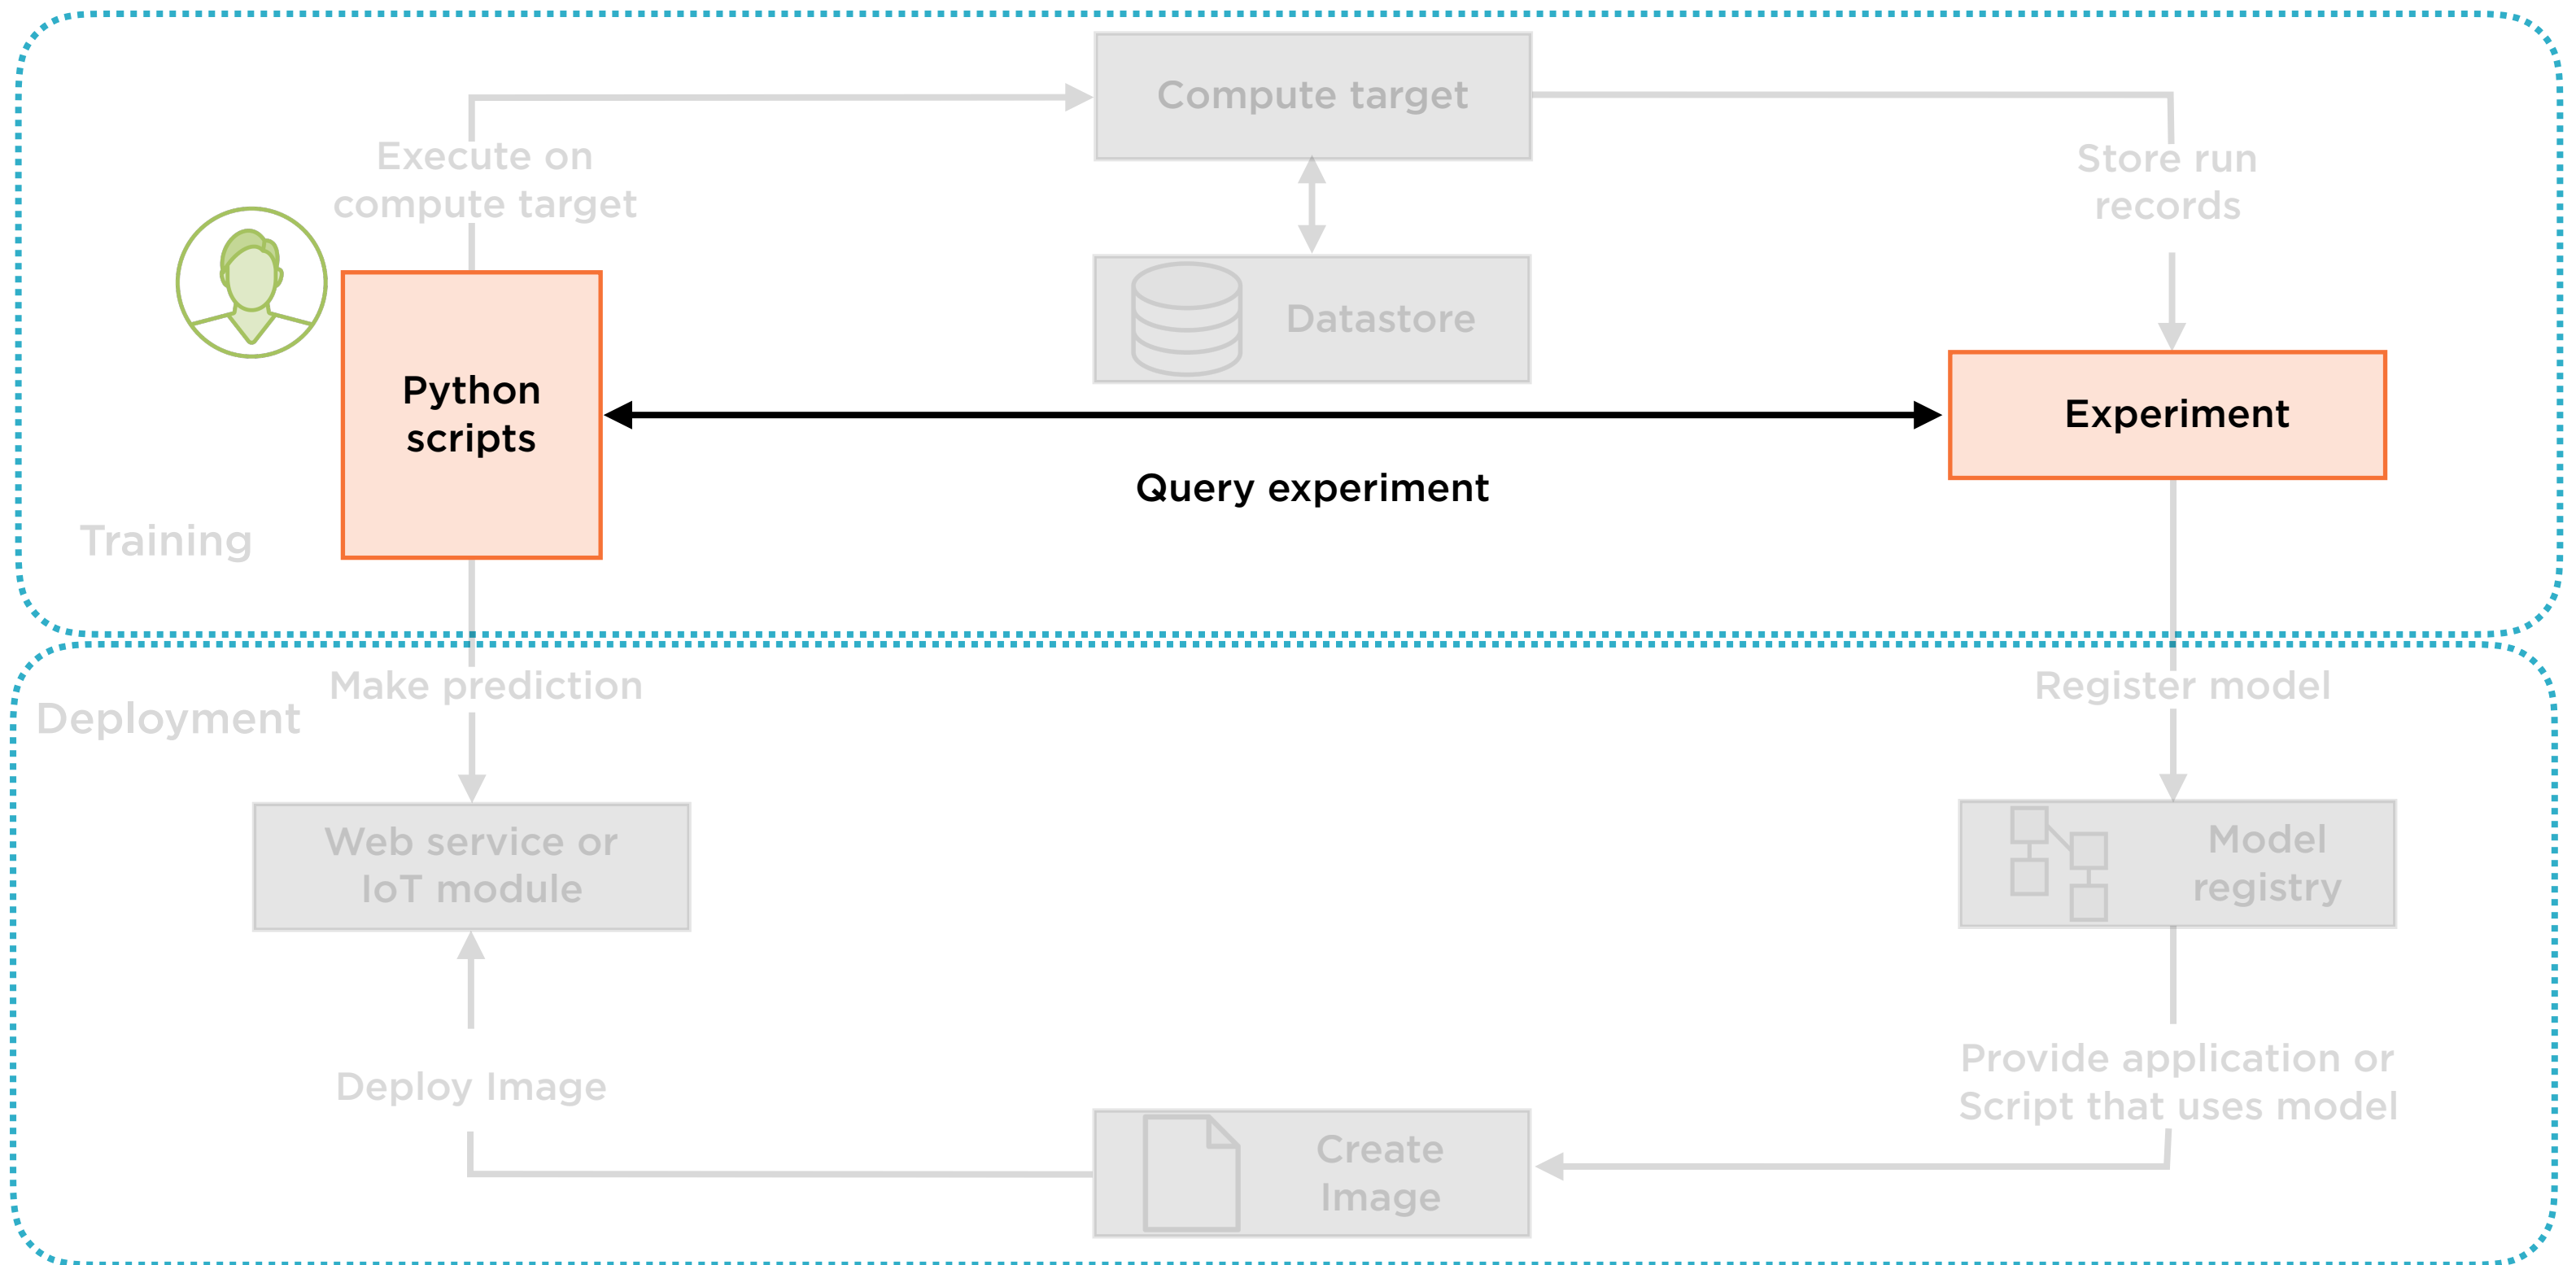
Storage Abstraction



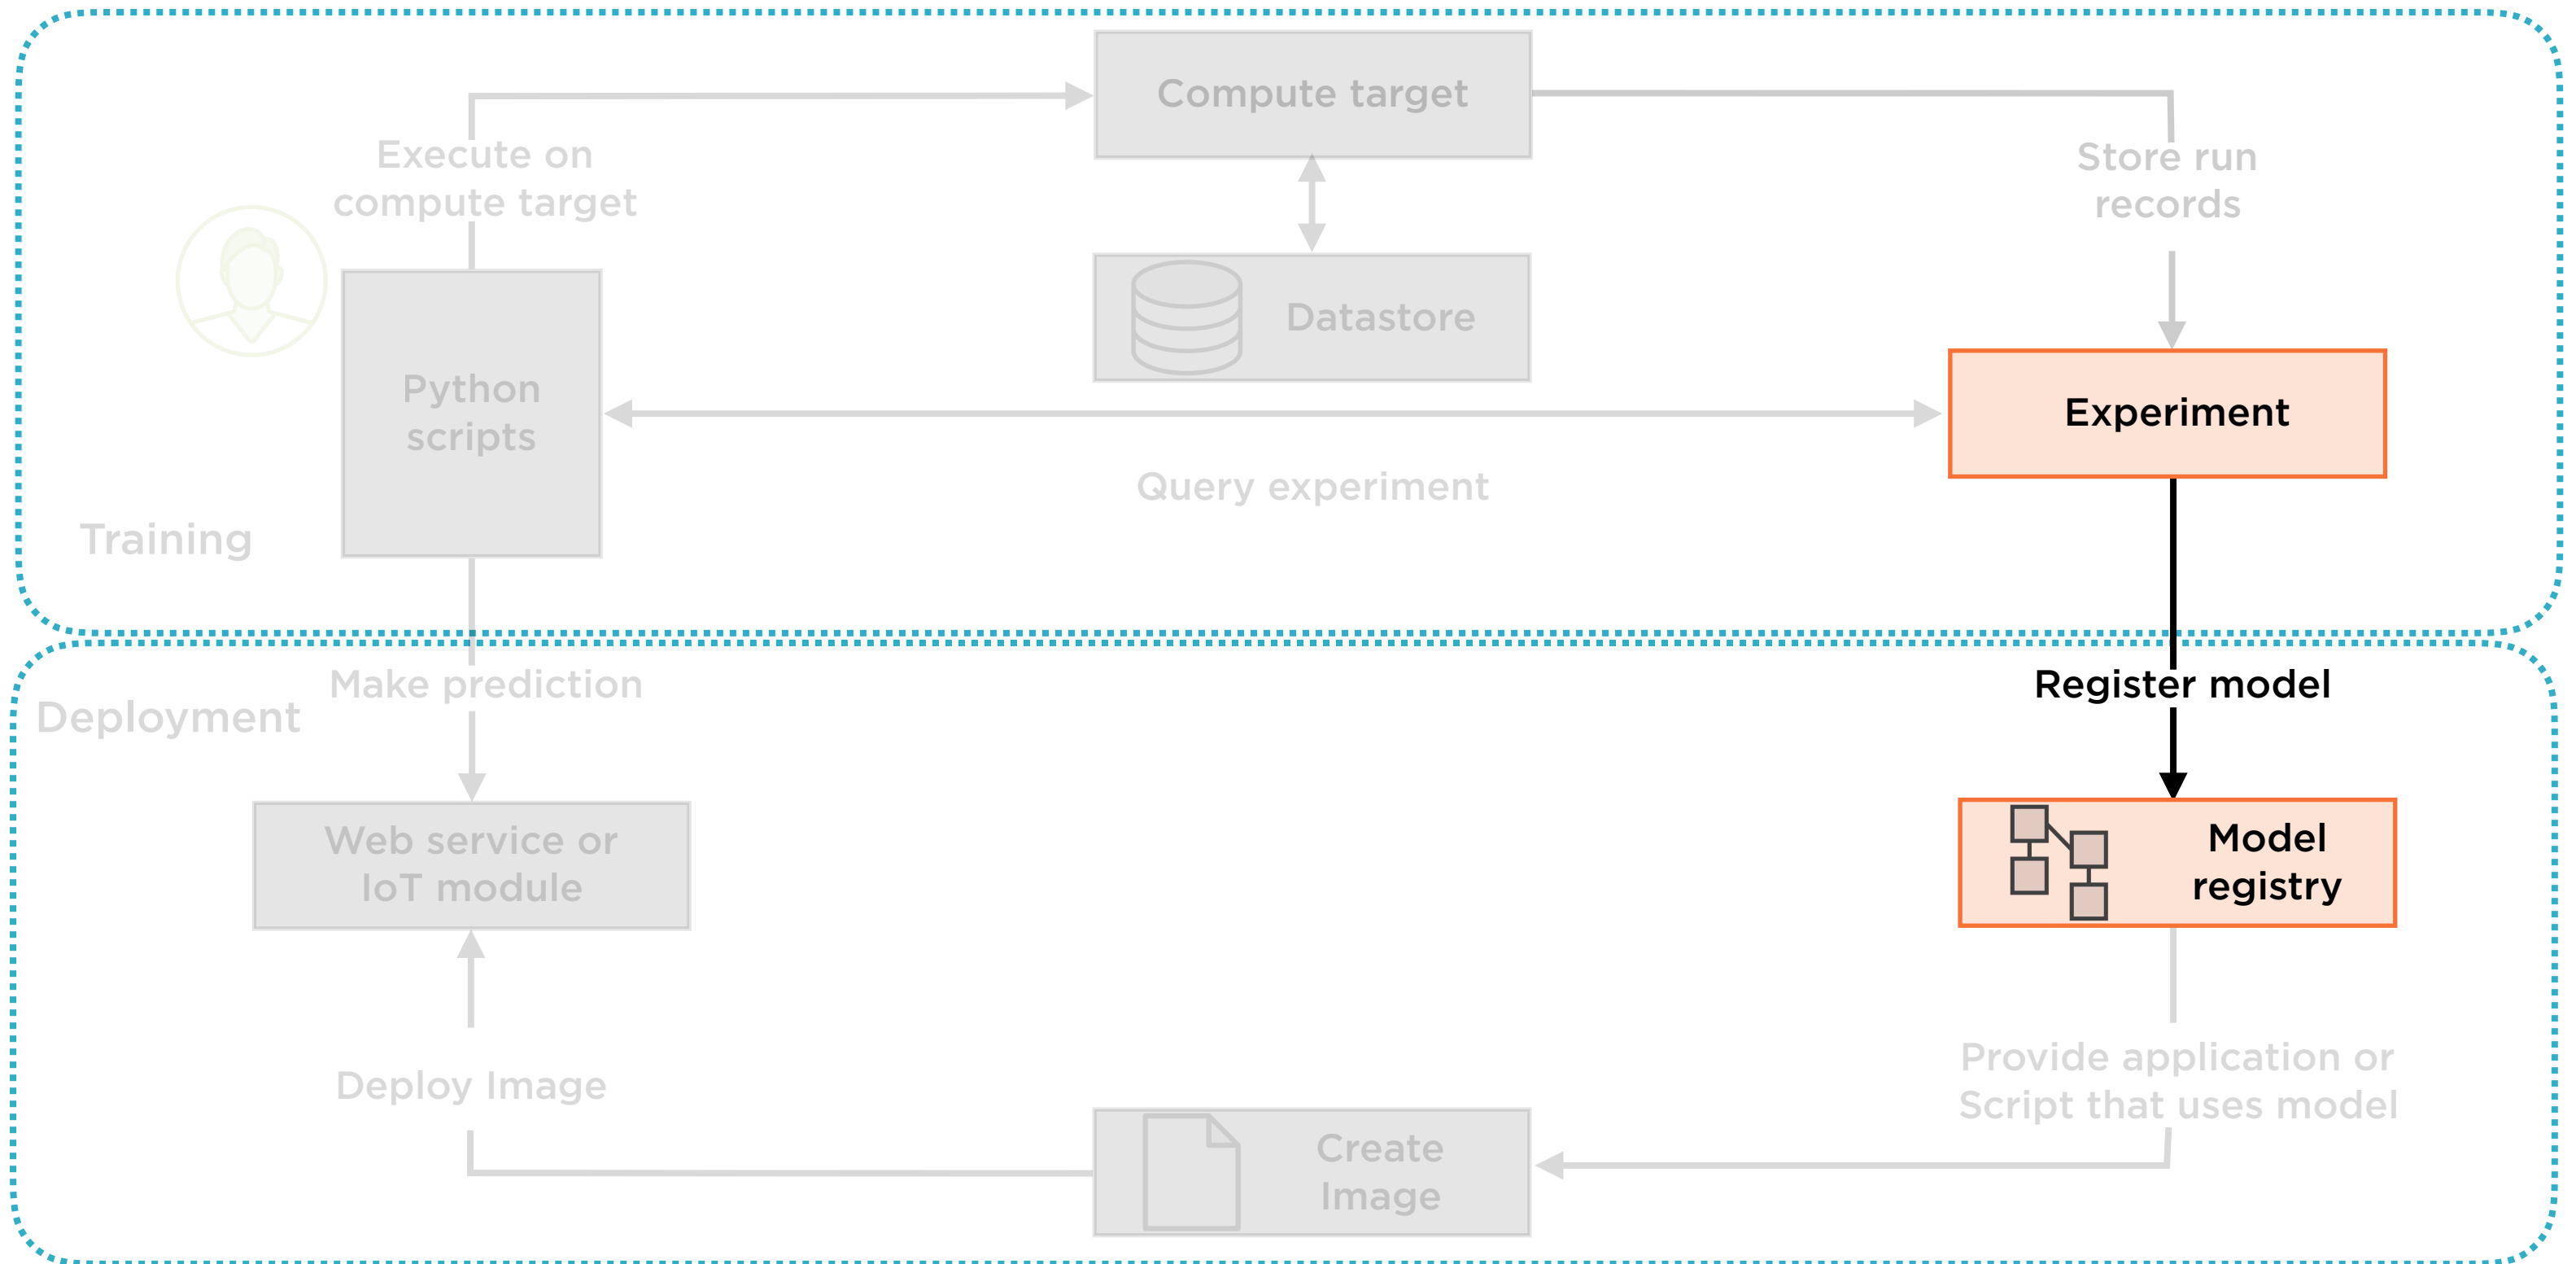
Groups Runs From the Same Script



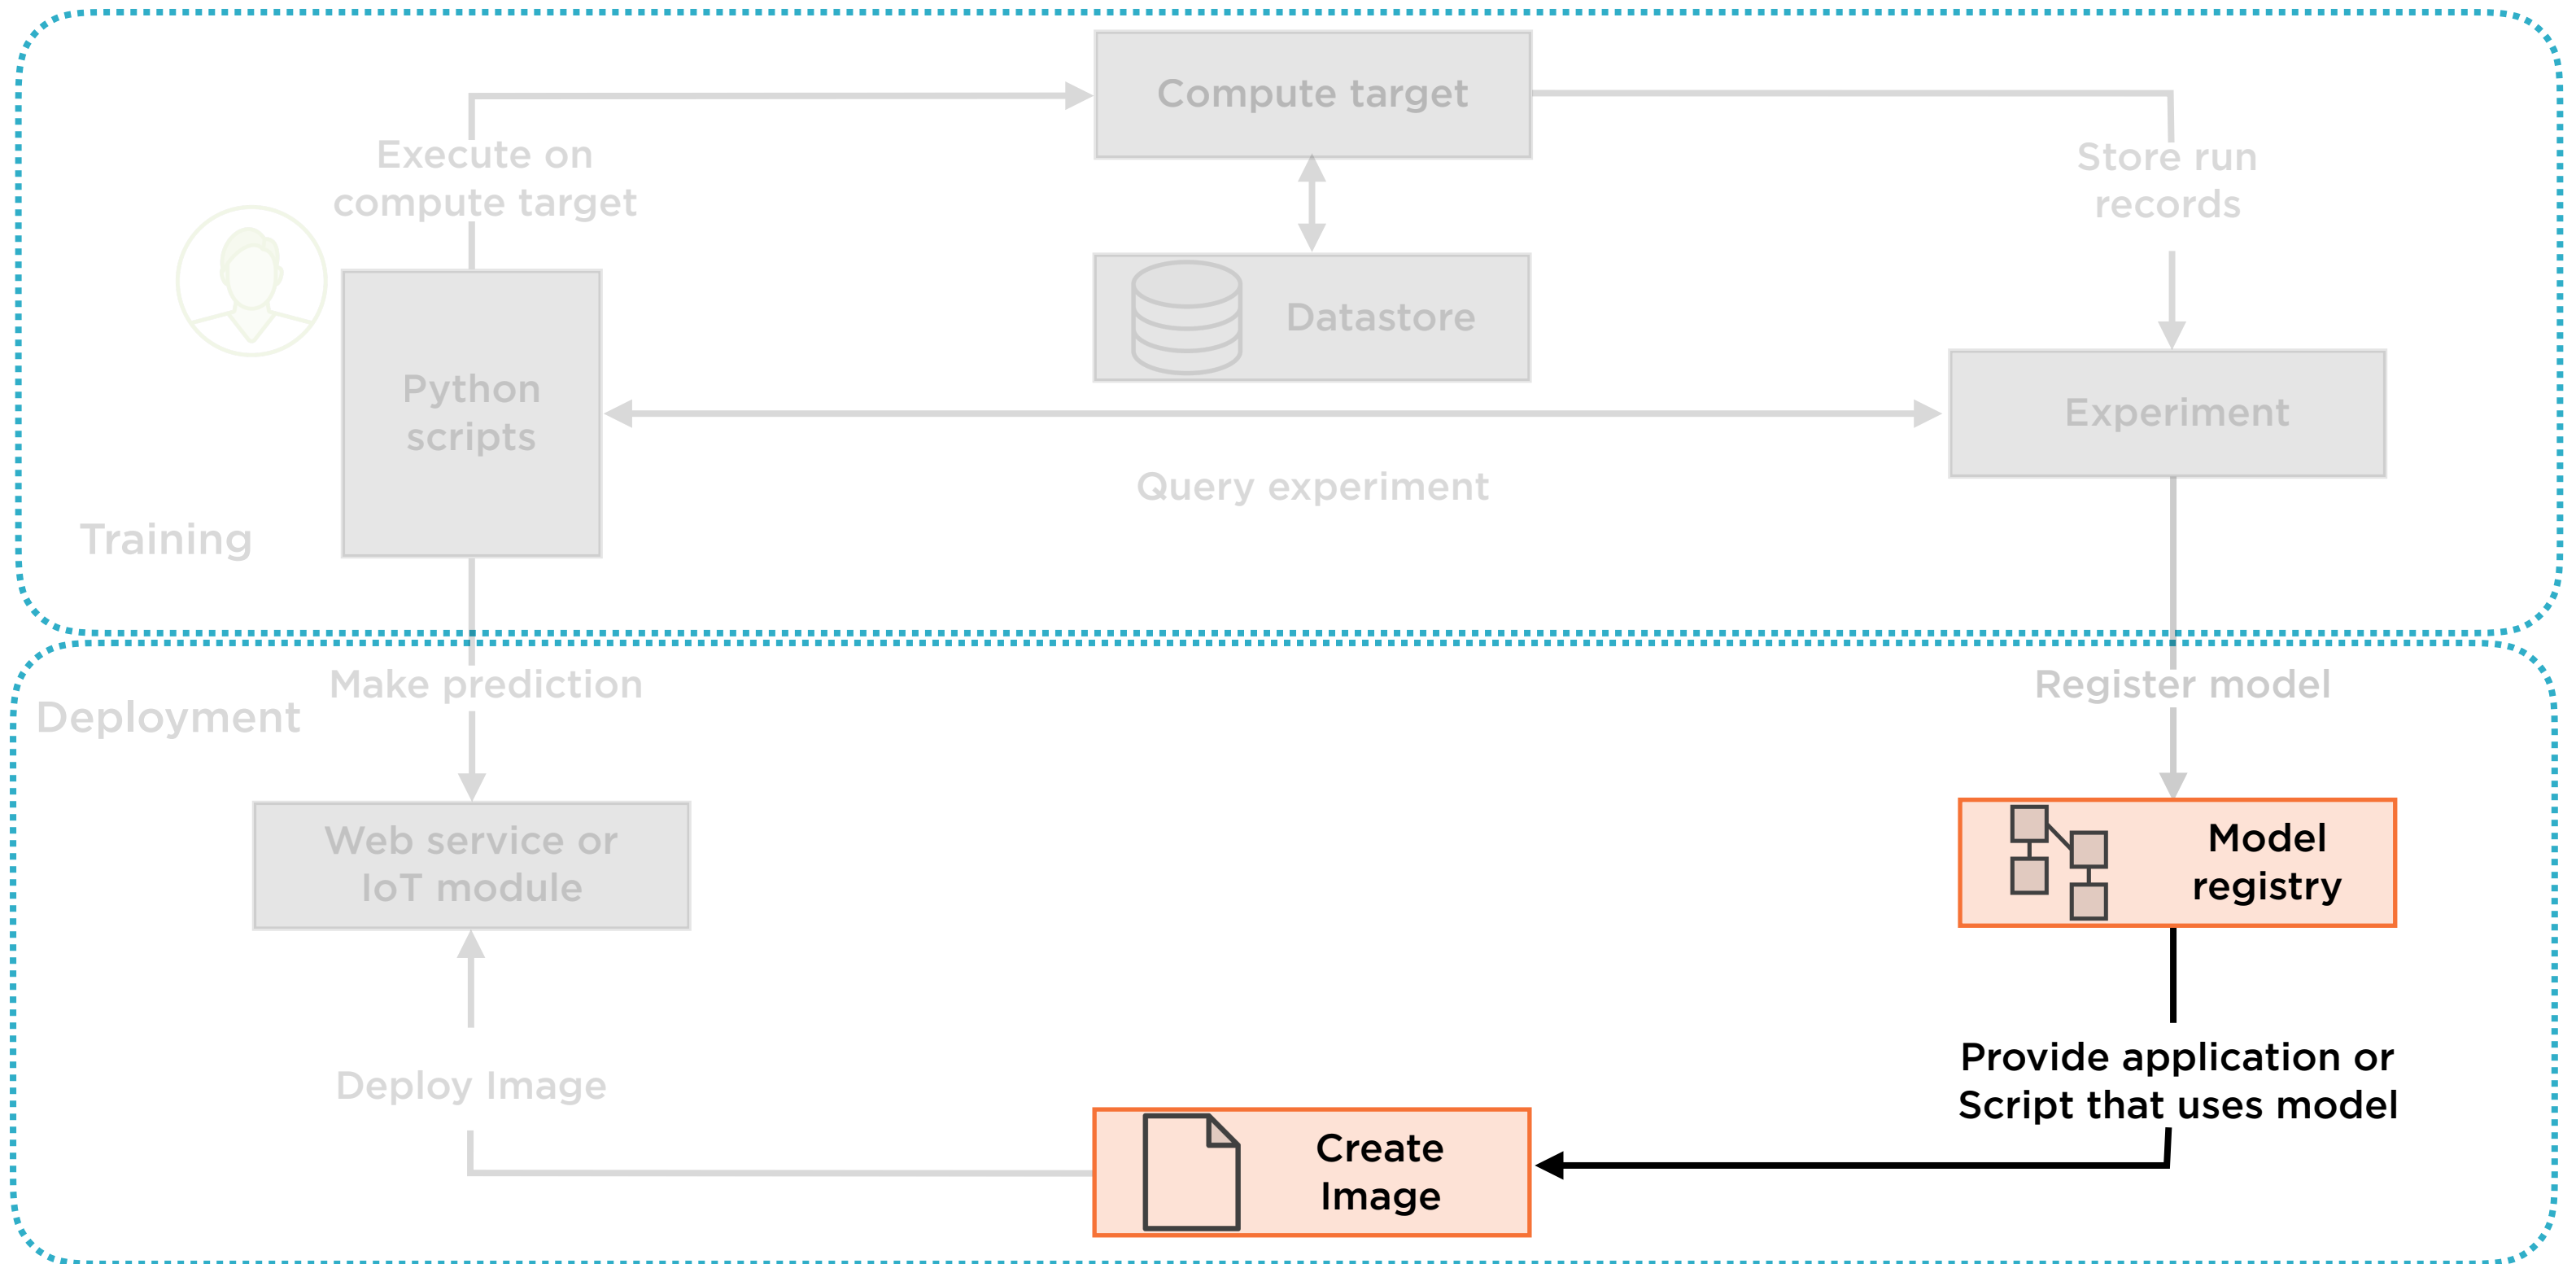
Groups Runs From the Same Script



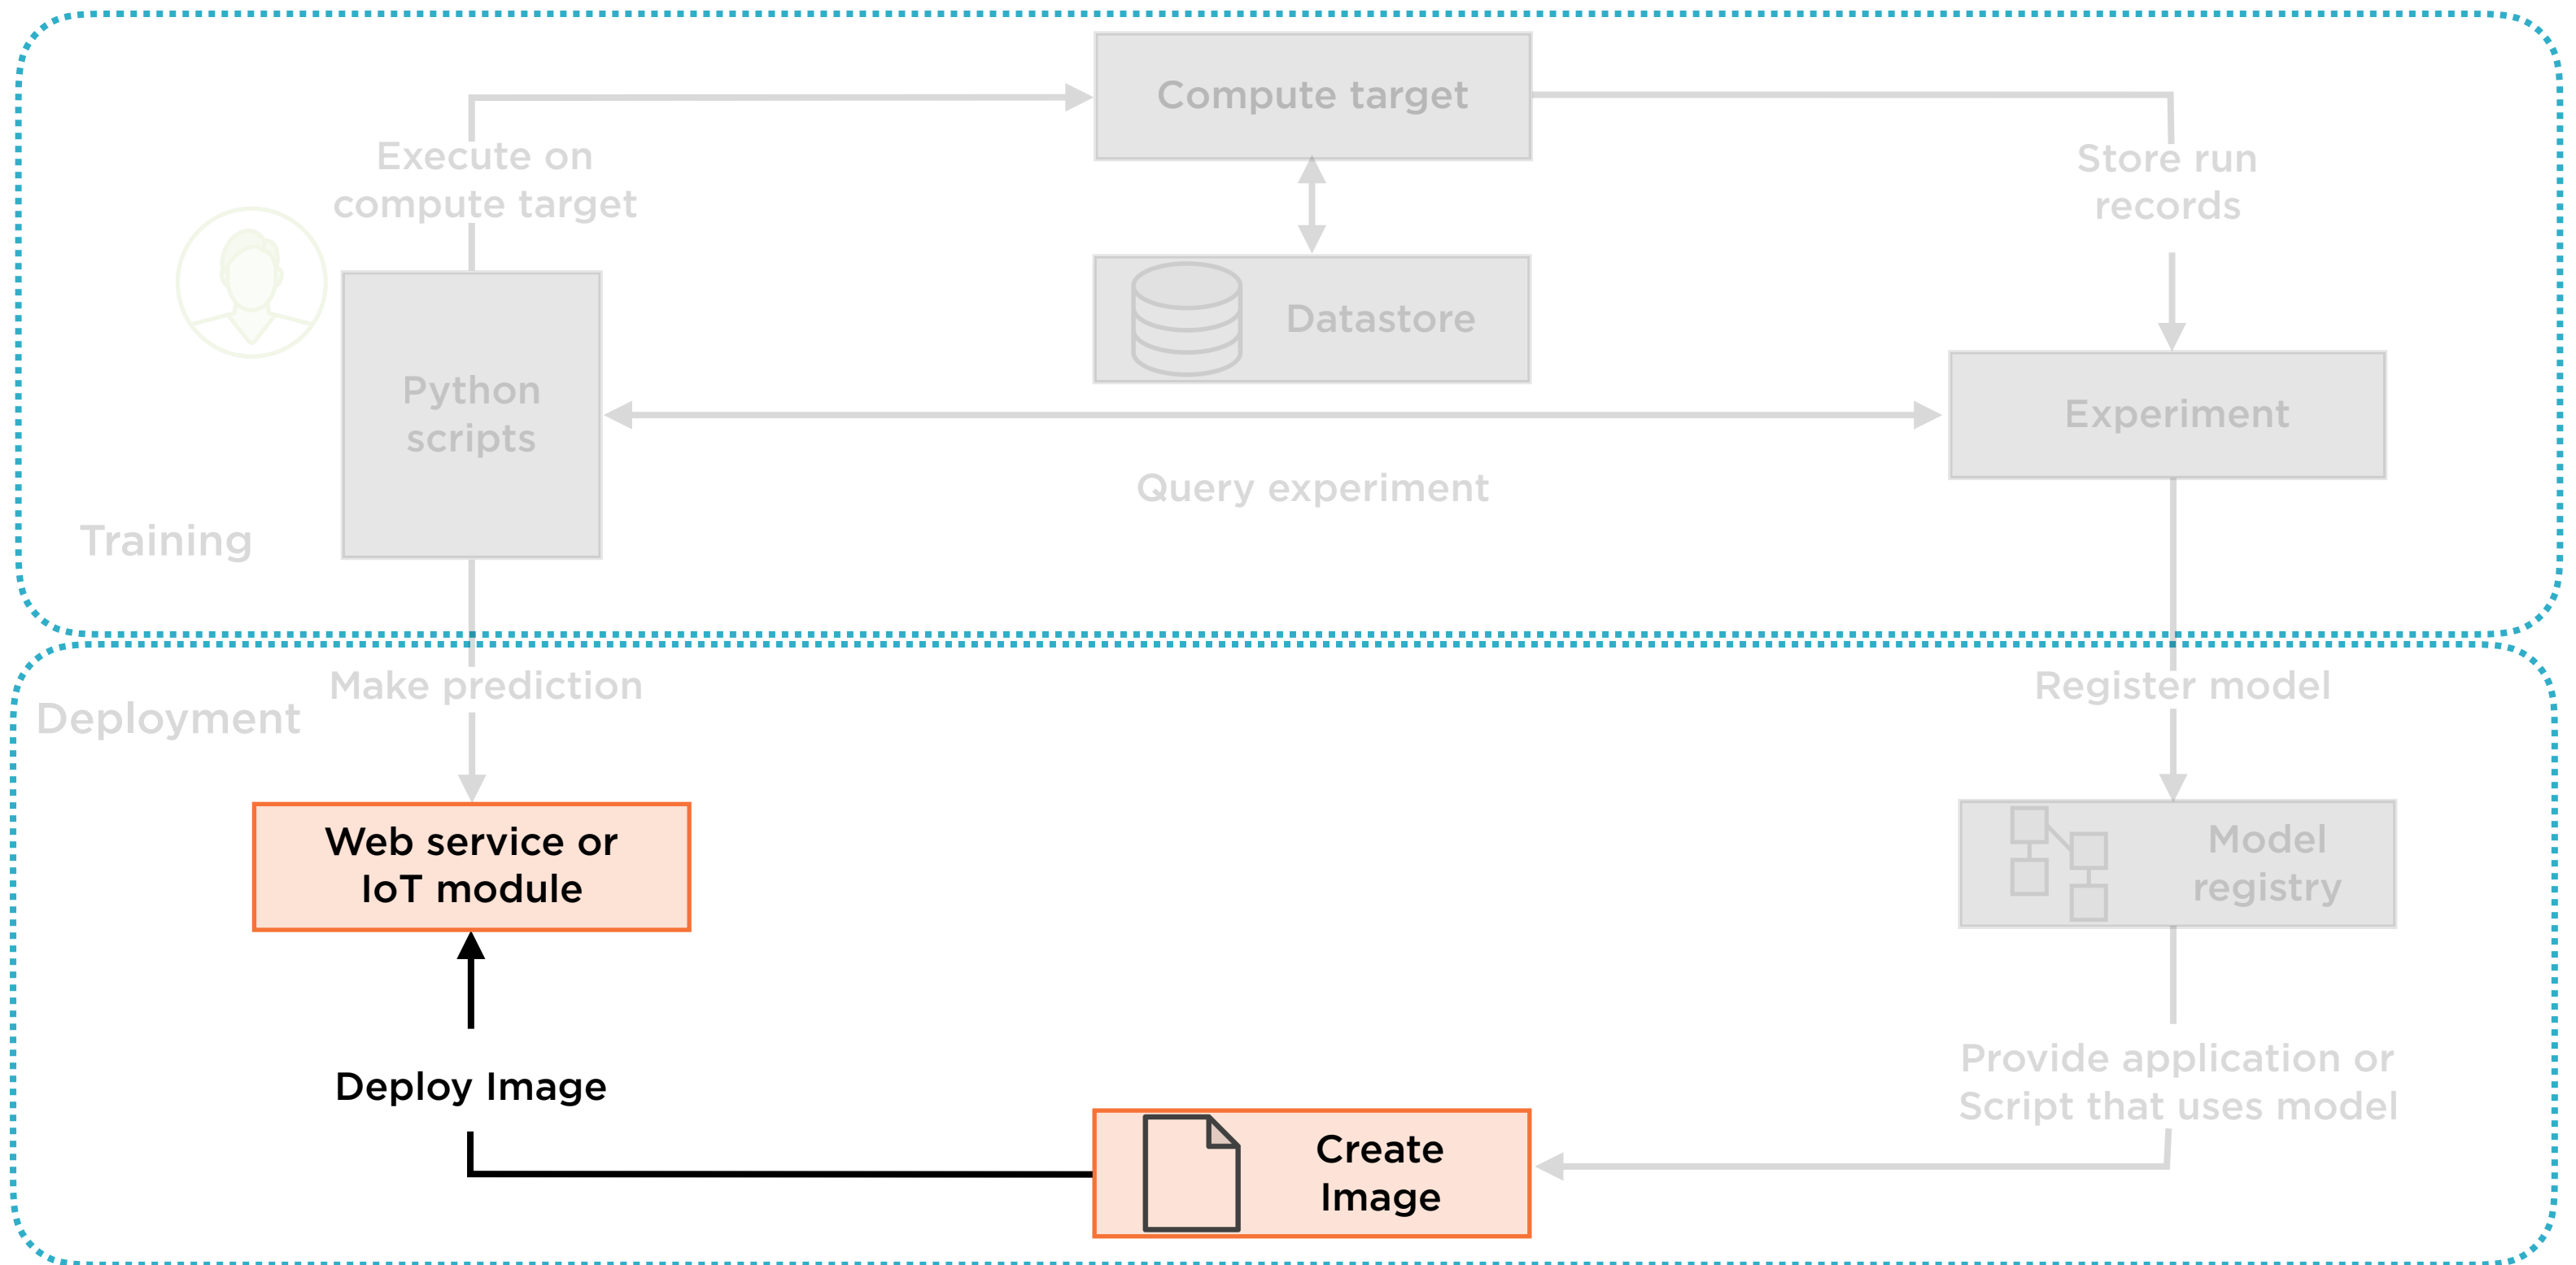
Azure Machine Learning Service



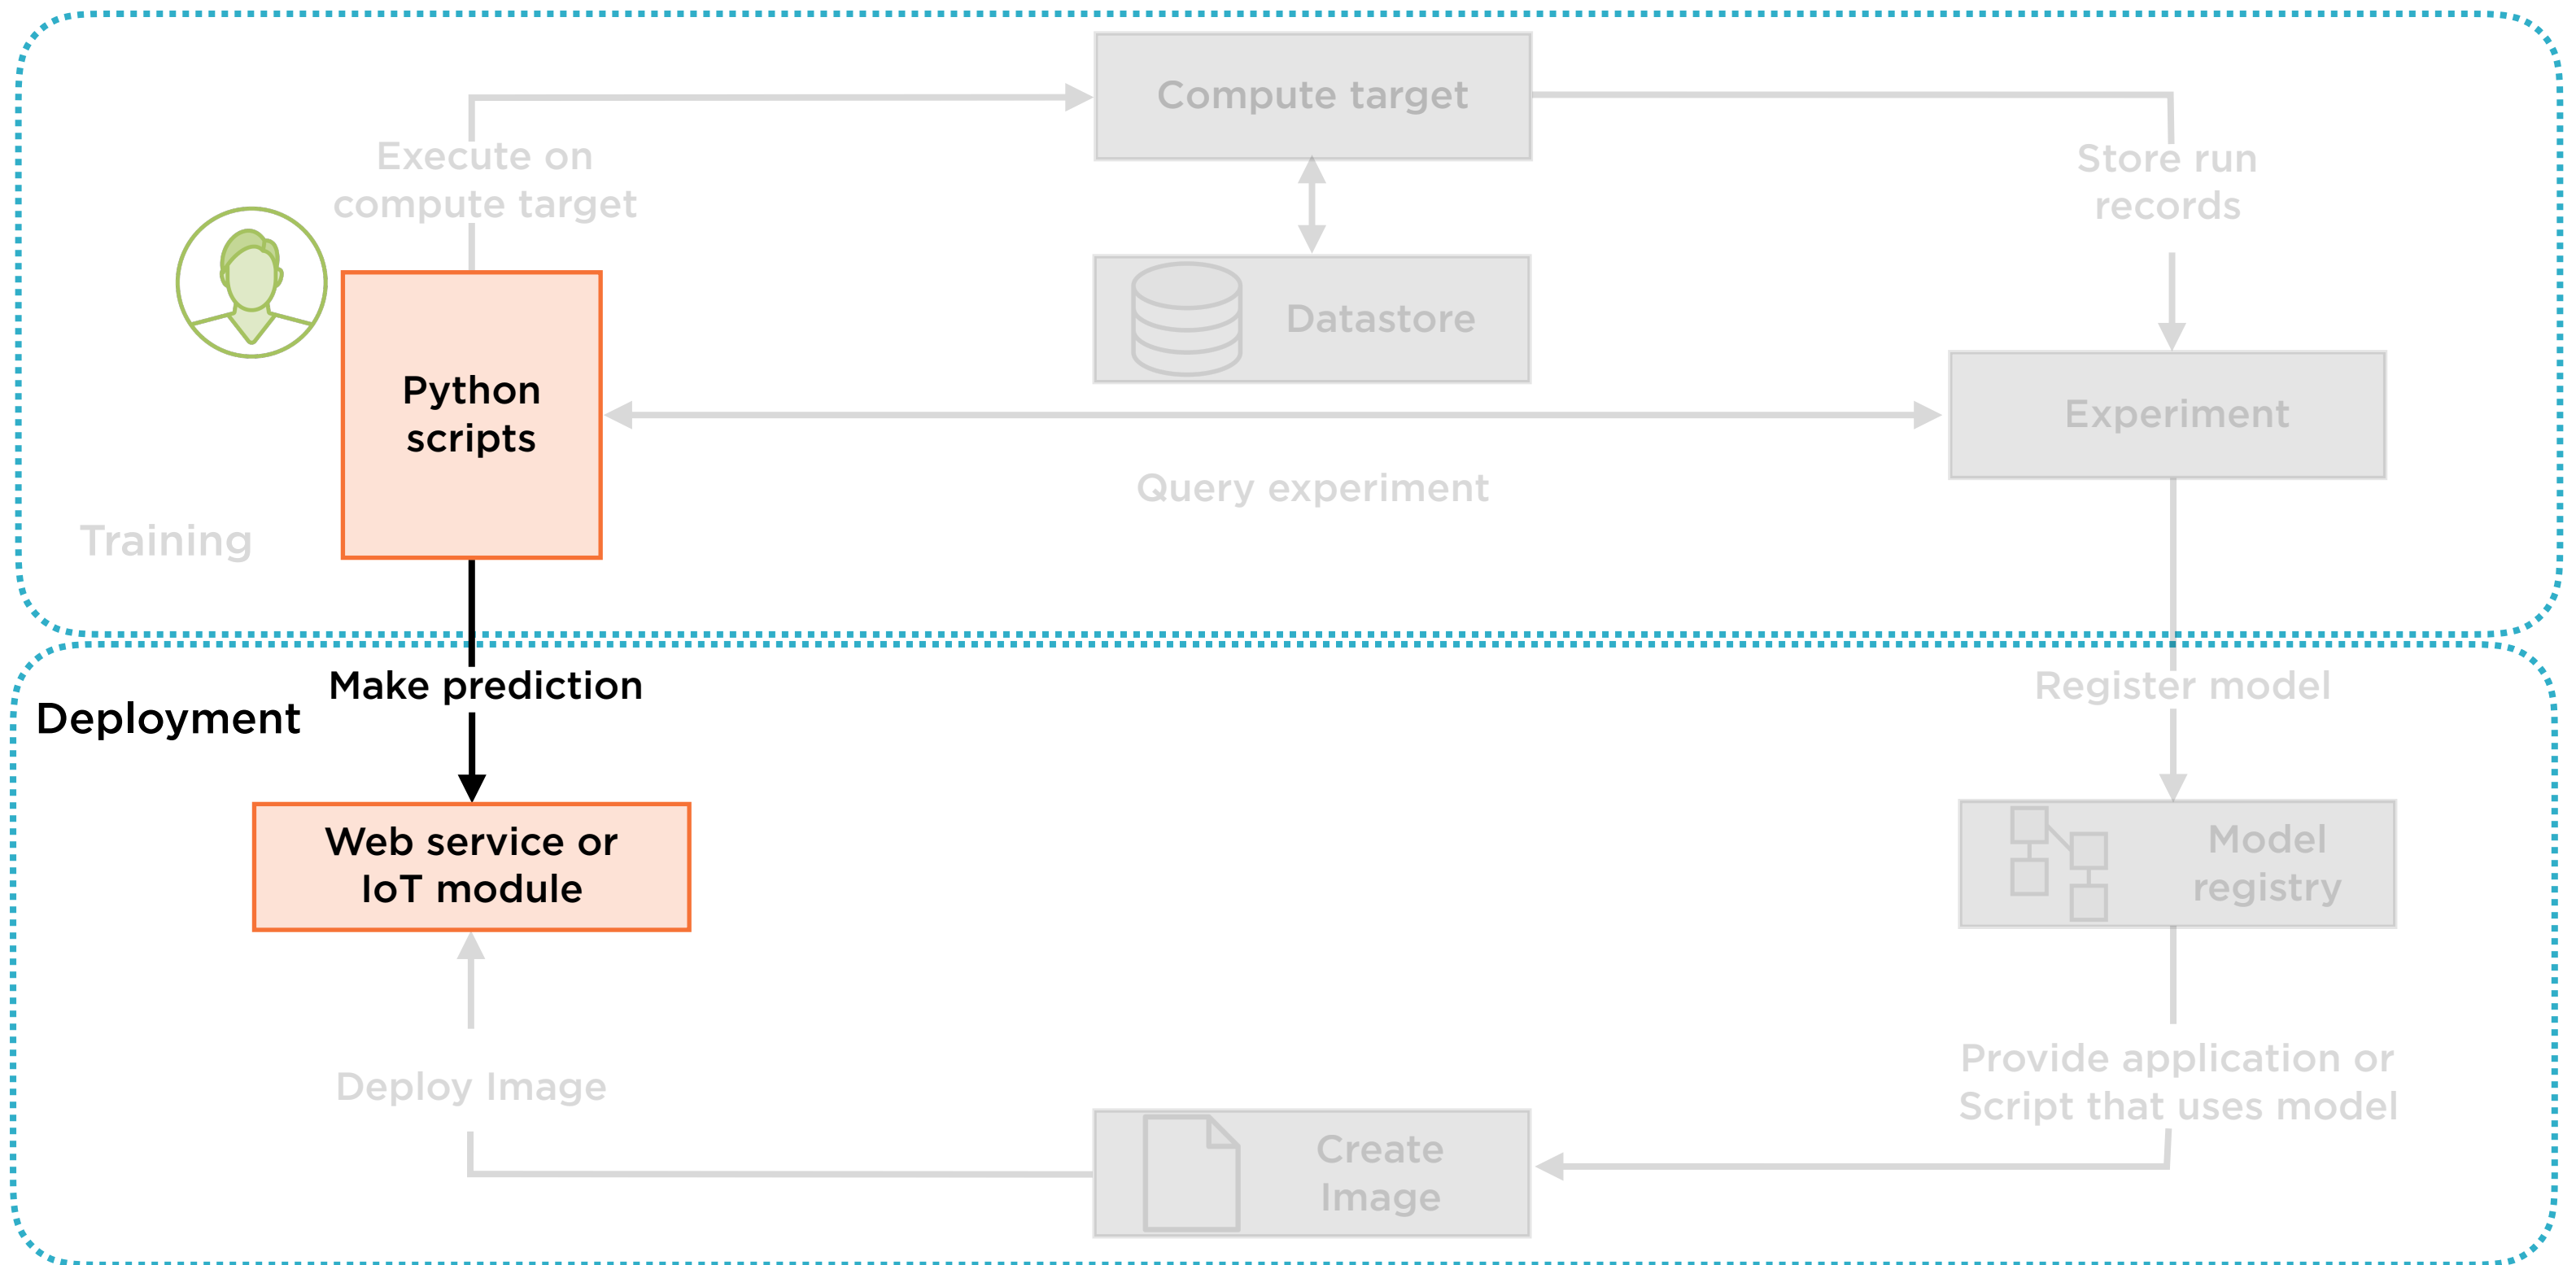
Azure Machine Learning Service



Azure Machine Learning Service



Azure Machine Learning Service



PyTorch Estimators on Azure



PyTorch on Azure Machine Learning



Azure Machine Learning Python SDK

PyTorch, TensorFlow, Chainer estimators

Available for different compute targets

Single Azure VM

GPU Cluster



Workspace

Top-level resource in Azure ML service, with list of compute targets, training runs, logs, metrics, and scripts.




Model

In context of Azure ML Service, a trained piece of code that has been registered with the workspace. Training may have been on Azure ML Service, or elsewhere.



Workspaces in Azure ML Service



Register model with workspace

Add scoring script to create image

Deploy image

- Azure container instances
- Azure Kubernetes Service
- RESTful HTTP endpoints

Model Registry

Keeps track of all models registered in a given workspace; maintains model version numbers and metadata as well.




Compute Targets

Compute resources where training scripts are run or models are hosted.



Compute Targets for Training



Local computer

Azure Machine Learning compute


Deep Learning or Data Science VM

Azure Databricks

Azure Data Lake Analytics

Apache Spark for HDInsight

Compute Targets for Deployment



Azure container instances

Azure Kubernetes Service

Azure IoT Edge

HTTP endpoint on FPGA

- Field-Programmable Gate Array
- Project Brainwave

Image

Unit of model deployment. Includes model, scoring script, and dependencies. Azure currently supports Docker images and FPGA images.



Experiment

Grouping of many runs of a single script. Contained within a workspace.




Horovod for Distributed Deep Learning

Horovod

Open-source toolkit for distributed training of deep learning models



Horovod



Originated at Uber

Initially for TensorFlow

Now supports PyTorch and Keras too



Horovod




Distributed training on raw TF is notoriously hard

Solutions such as GCP ML Engine are platform-specific

Horovod serves as platform-agnostic solution



Horovod



Gained popularity for ease-of-use as well as speed

Great for large-scale GPU training



Demo

**Running distributed training in PyTorch
with Azure Estimators and the Horovod
framework**



Azure Deep Learning VM



Azure Deep Learning VM



Specially configured VM with GPU support

Pre-provided with samples of image and text processing models

Variant of Azure Data Science VM

- Pre-configured with development and visualization tools



Azure Deep Learning VM



Vertical scaling - one powerful VM

Runs on Azure GPU NC-series VMs

Optimized to almost bare-metal

Variant of Data Science VM

- Windows 2016
- Ubuntu



Azure Deep Learning VM



Microsoft Cognitive Toolkit (CNTK)

TensorFlow

Keras

Caffe, Caffe2

MXNet

Torch, PyTorch



The data science and deep learning software is free - you only pay for the underlying VM



Demo

**Building PyTorch models on the Azure
Deep Learning VM**



Summary

Azure Machine Learning Service

Azure notebooks for PyTorch

PyTorch estimators on Azure

**Leverage Horovod for distributed
deep learning**

Azure Deep Learning VMs

